

Discussion Worksheet 6

Note that the code generation conventions used throughout this worksheet are for the “small language” described in lecture and not for ChocoPy.

1 Code Generation

A stack machine is an evaluation model where all data is stored on the stack. There are no variables or registers. Every operation in this model does the following:

1. Pop some number of arguments from the top of the stack.
2. Compute something over these arguments.
3. Push a result onto the stack.

E.g., “add” pops two elements off the stack, computes their sum, and pushes the sum onto the stack.

We can optimize this model by adding an “accumulator” register, which usually contains the top element of the stack. We can then manipulate this register directly, instead of operating only in the stack (in memory).

Exercise 0 (warmup) Try filling in the model code for the pure stack machine and the accumulator stack machine, given the expression $(7 + 5)$:

	Pure Stack Machine	With Accumulator
Store 7	push 7	acc = 7
	—	
Store 5	—	—
	—	
Compute sum	—	

	Pure Stack Machine	With Accumulator
Store 7	push 7	acc = 7
	—	
Store 5	push 5	push acc; acc = 5
	—	—
Compute sum	add	acc += top of stack; pop
	—	

Note that we may implement the bottom-right cell (computing the sum, in the accumulator-stack model) in RISC-V as follows, if the accumulator is stored in `a0`:

```

t1 <- top    # load top of stack into temporary register
add a0, t1, a0 # perform the addition
pop          # pop the stack

```

We will implement this accumulator stack machine in machine code using RISC-V. Here are the code generation conventions we will use, based on those we used for the “small language” in lecture.

1. We keep the accumulator in register `a0`.
2. We use `t1` to load elements from the stack for operations, since we cannot operate directly on the stack.
3. The next writeable stack location is at address `sp`. When we write (push), we decrease `sp` by 4.

4. The last readable stack location is at address $sp + 4$.

For example, here is the `cgen` function for a few different constructs:

<code>cgen(e1 + e2) =</code>	<code>cgen(f(e1,...,en)) =</code>	<code>cgen(def f(x1,...,xn) = e) =</code>
<code> cgen(e1)</code>	<code> push fp</code>	<code>f_entry:</code>
<code> push a0</code>	<code> cgen(en)</code>	<code> mv fp, sp</code>
<code> cgen(e2)</code>	<code> push a0</code>	<code> push ra</code>
<code> t1 <- top</code>	<code> ...</code>	<code> cgen(e)</code>
<code> add a0, t1, a0</code>	<code> cgen(e1)</code>	<code> ra <- top</code>
<code> pop</code>	<code> push a0</code>	<code> addi sp, sp, z</code>
	<code> jal f_entry</code>	<code> lw fp, 0(sp)</code>
		<code> jr ra</code>

Exercise 1 Let's convert the following pseudocode into RISC-V using the accumulator model:

```
factorial 0 = 1
factorial x = x * factorial (x - 1)
```

1.1 What elements will be in each Activation Record for each call to factorial?

old FP (frame pointer), x (argument), and RA (return address).

1.2 Fill in the RISC-V code for the factorial function:

```
factorial_entry:
#entry: set up our FP, and save the return address for later
-                               #Set new FP
-                               #Store RA
#body
-                               #acc = x
-                               #push acc
-                               #acc = 0
-                               #load top of stack to temporary
-                               #pop to maintain stack hygiene
-                               #jump to right branch
false_branch:
-                               #acc = x
-                               #push acc
-                               #start evaluation of (factorial(x-1))
-                               #acc = x
-                               #push acc
-                               #acc = 1
-                               #load top of stack to temporary
-                               #perform the subtraction (x - 1)
-                               #pop to maintain stack hygiene
-                               #push the last argument, (x-1)
-                               #recursively call 'factorial'
-                               #start computing the product (x * factorial(x - 1))
-                               #perform the multiplication
-                               #maintain stack hygiene
-                               #skip over the true_branch
true_branch:
-                               #compute (0)
end_if:
#exit
-                               #Load RA which we saved earlier
addi sp,sp,z                    #Jump to old FP. What is z? z=___
-                               #Set FP to be the old FP, which the callee saved on the stack
-                               #Return to caller
```

You may use the *cgen* function described in lecture to find this code. Try it and see if you get the same assembly code as we did! (Note that “#...” is a comment.)

```
factorial_entry:
#entry: set up our FP, and save the return address for later
    mv fp, sp                #Set new FP
    push ra                  #Store RA
#function body
#first we evaluate (x == 0), and then branch
    lw a0, 4(fp)             #acc = x
```

```

    push a0                #push acc
    li a0, 0               #acc = 0
    t1 <- top              #load top of stack to temporary
    pop                   #pop to maintain stack hygiene
    beq a0, t1, true_branch #jump to right branch
false_branch:
#code to evaluate (x * factorial(x - 1))
#first we evaluate (x), and save the result to the stack
#then we evaluate (factorial(x - 1))
#and then multiply the two values
    lw a0, 4(fp)          #acc = x
    push a0               #push acc
    push fp               #start evaluating (factorial(x-1))
                           #evaluate the last argument, (x-1), in the next 6 lines
    lw a0, 4(fp)          #acc = x
    push a0               #push acc
    li a0, 1              #acc = 1
    t1 <- top             #load top of stack to temporary
    sub a0, t1, a0         #perform the subtraction
    pop                   #pop to maintain stack hygiene
    push a0               #push the last argument, (x-1)
    jal factorial_entry    #recursively call 'factorial'
    t1 <- top             #result is in a0; now compute the product
    mul a0, t1, a0         #a0 is now (x * factorial(x-1))
    pop                   #maintain stack hygiene
    j end_if
true_branch:
#code to evaluate (0)
    li a0, 1
end_if:
#exit: load our RA from earlier, and restore the SP and FP
    ra <- top             #Load RA which we saved earlier
    addi sp,sp,z          #Jump to old FP. z=8+4(num_args)=12
    lw fp, 0(sp)          #Set FP to be the old FP, which the callee saved on the stack
    jr ra                 #Return to caller

```

There are certainly a lot of opportunities for optimization! Can you spot any ways to make the code shorter or more efficient?

Here is an iterative definition for factorial:

```
def factorial(x):
    total = 1
    while (x > 0):
        total = total * x
        x = x - 1
    return total
```

1.3 Fill in the RISC-V code for the iterative factorial. Again, use the code generation conventions for the small language described in lecture, as outlined on page 1 of the worksheet.

Note that local variables also are stored in the stack, like function arguments. Let's put any local variables on the stack just below the RA (i.e. lower addresses); they can be accessed using a fixed offset from `fp`, just like any parameters. Which offset(s) should we use?

```
factorial_entry:
    -
    -                                # set up our FP
    -                                # save RA to the stack
#body
    -                                # evaluate the initial value for 'total'
    -                                # push 'total' to the stack, for storage throughout the
    -                                # execution of this function
while_loop:
    -                                # evaluate (x > 0) and branch in the next few lines
    -
    -
    -
    -
    -                                # branch as necessary. hint: 'bge' = branch if greater or equal
    -                                # evaluate (total * x) in the next few lines
    -
    -
    -
    -
    -                                # store accumulator in 'total'
    -                                # evaluate (x - 1) in the next few lines
    -
    -
    -
    -
    -                                # store accumulator in 'x'
    -                                # continue to the next iteration
    j while_loop
exit_while:
#exit
    -                                # load return value into a0
    -                                # maintain stack hygiene: pop our storage for 'total'
    -                                # load return address
    addi sp,sp,z                    # restore stack. z=___?
    -                                # restore old FP
    -                                # return to caller
```

factorial_entry:

```

        mv fp, sp
        push ra
#body
        li a0, 1                # evaluate and save 'total'; note that it can be accessed at -4(fp)
        push a0
        # fall through to the start of the 'while' loop
while_loop:
        lw a0, 4(fp)            # evaluate (x > 0) in the next few lines
        push a0
        li a0, 0
        t1 <- top
        pop
        bge a0, t1, exit_while  # bge = branch if greater than or equal
        lw a0, -4(fp)           # evaluate (total * x) in the next few lines
        push a0
        lw a0, 4(fp)
        t1 <- top
        mul a0, t1, a0
        pop
        sw a0, -4(fp)           # now a0 = (total * x); store it in 'total'
        lw a0, 4(fp)           # evaluate (x - 1) in the next few lines
        push a0
        li a0, 1
        t1 -> top
        sub a0, t1, a0
        pop
        sw a0, 4(fp)           # now a0 = (x - 1); store it in 'x'
        j while_loop           # continue to the next iteration
exit_while:
#exit
        lw a0, -4(fp)           # set return value to 'total'
        pop                    # maintain stack hygiene: we previously 'push'ed 'total' to
                                # the stack, so now we 'pop' it
        ra <- top               # standard function epilogue from here to the end
        addi sp, sp, 12
        lw fp, 0(sp)
        jr ra

```

2 Temporaries

Instead of pushing and popping elements from the stack, it is sometimes more efficient to precompute locations for temporary variables, allocate stack space at the beginning, and then access them directly. A function can perform the preallocation before evaluating its body expression, by checking $NT(expr)$ and allocating that much space on the stack.

Let's consider the expression $3 + (5 + 7) + 9$. This requires two temporaries. (What change could we make to $(5+7)$ to make it require three temporaries?) Therefore, we will preallocate two spaces on the stack, accessible via $-4(fp)$ and $-8(fp)$.

Exercise 2 Recall that $cgen(e_1 + e_2, nt)$ is defined as

```
cgen(e1 + e2, nt) =
    cgen(e1, nt)
    sw a0, -nt(fp)
    cgen(e2, nt + 4)
    lw t1, -nt(fp)
    add a0, t1, a0
```

2.1 Fill out the blanks in the “Preallocate” column, so that they are the equivalents of the instructions in the “Push/Pop” column, but using fp -relative addressing instead of **push** and **pop**. What should each offset from fp be? How can we compute this ahead of time?

Push/Pop	Preallocate	$cgen(3+(5+7)+9, nt=4)$
=====	=====	$cgen(3+(5+7), nt=4)$
$a0 = 3$	$a0 = 3$	$cgen(3, nt=4)$
push $a0$	-	Save temp
=====	=====	$cgen(5+7, nt=8)$
$a0 = 5$	$a0 = 5$	$cgen(5, nt=8)$
push $a0$	-	Save temp
$a0 = 7$	$a0 = 7$	$cgen(7, nt=12)$
pop $t1$	-	Load temp 5
add $a0, t1, a0$	add $a0, t1, a0$	Compute $5+7$
pop $t1$	-	Load temp 3
add $a0, t1, a0$	add $a0, t1, a0$	Compute $3+(5+7)$
push $a0$	-	Save temp
$a0 = 9$	$a0 = 9$	$cgen(9, nt=8)$
pop $t1$	-	Load temp $3+(5+7)$
add $a0, t1, a0$	add $a0, t1, a0$	Compute full sum

Push/Pop	Preallocate	$cgen(3+(5+7)+9, nt=4)$
$a0 = 3$	$a0 = 3$	$cgen(3+(5+7), nt=4)$
push $a0$	sw $a0, -4(fp)$	$cgen(3, nt=4)$
		Save temp
$a0 = 5$	$a0 = 5$	$cgen(5+7, nt=8)$
push $a0$	sw $a0, -8(fp)$	$cgen(5, nt=8)$
$a0 = 7$	$a0 = 7$	Save temp
pop $t1$	lw $t1, -8(fp)$	$cgen(7, nt=12)$
add $a0, t1, a0$	add $a0, t1, a0$	Load temp 5
pop $t1$	lw $t1, -4(fp)$	Compute $5+7$
add $a0, t1, a0$	add $a0, t1, a0$	Load temp 3
push $a0$	sw $a0, -4(fp)$	Compute $3+(5+7)$
$a0 = 9$	$a0 = 9$	Save temp
pop $t1$	lw $t1, -4(fp)$	$cgen(9, nt=8)$
add $a0, t1, a0$	add $a0, t1, a0$	Load temp $3+(5+7)$
		Compute full sum

2.2 Why isn't it a problem that we have 'cgen(7, nt=12)' even though we only allocated 8 bytes on the stack?

'NT("7") = 0', so no temporaries are stored when evaluating this expression.